

Regression and SVM

October 27, 2015

Abstract

1 Logistic Regression

In past discussions, we've often encountered the objective function of our MLE estimate, or of some other estimate that we've established for some parameter θ that we'd like to minimize so as to minimize the value of our error function. This is called **regression**. Here we will discuss **logistic regression**, which involves fitting the data to a curve of the form

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

This is known as the **sigmoid** function, and it's often used when the data we would like to fit is well-modeled as a classification problem (with labels +1 and -1). Given d -dimensional data $x^{(1)}, x^{(2)}, \dots, x^{(n)}$ and corresponding scalar labels $y^{(1)}, y^{(2)}, \dots, y^{(n)}$, we can define our LR objection function as:

$$\text{NLL}(w) = \sum_{i=0}^n \log(1 + e^{-y^{(i)}(x^{(i)} \cdot w + w_0)})$$

given the parameters w, w_0 . To reduce overfitting, we can apply an $L2$ regularization term to our equation, as in ridge regression. As a result, we are simply minimizing

$$E_{LR}(w) = \text{NLL}(w) + \lambda w^T w$$

so that

$$w^* = \operatorname{argmin}_w \left[\sum_{i=0}^n \log(1 + e^{-y^{(i)}(x^{(i)} \cdot w + w_0)}) + \lambda w^T w \right]$$

2 Support Vector Machines

Support Vector Machines are a machine-learning technique that is used to classify binary sets of data. We define θ as the normal to the decision hyperplane, and classify data points by $\operatorname{sgn}(\theta^T x + \theta_0)$. In order to train the classifier, we initially set up the minimization problem (known as **Hard-SVM**):

$$\max_{\theta, \theta_0} \frac{1}{\|\theta\|} \min_{1 \leq i \leq n} y^{(i)}(\theta^T x^{(i)} + \theta_0) \quad (1)$$

However, we cannot satisfy this minimization if the training set is not linearly separable. Therefore, we add a "slack variable" to each constraint, which is a measure of the "wrongness" of the decision boundary when classifying that point. We call these slack variables ξ_i . We then desire to minimize

$$\min_{\theta, \theta_0, \xi} \frac{\lambda}{2} \|\theta\|^2 + \frac{1}{n} \sum_{i=1}^n \xi_i, \quad (2)$$

$$s.t. \quad y^{(i)}(\theta^T x^{(i)} + \theta_0) \geq 1 - \xi_i, \quad (3)$$

$$\xi_i \geq 0, i \in n \quad (4)$$

However, we find it more computationally efficient to solve the dual form of the **Soft-SVM**:

$$\max_{\alpha \in \mathbb{R}^n} \left[\sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} (x^{(i)})^\top x^{(j)} \right] \quad (5)$$

$$s.t. \quad 0 \leq \alpha_i \leq C \quad (6)$$

$$\sum_i \alpha_i y^{(i)} = 0 \quad (7)$$

To provide some intuition: α_i is the weight of each training point on the final decision boundary. It is 0 for all $x^{(i)}$ that are farther from the decision boundary than the margin, so that only the “important” $x^{(i)}$ are “support vectors”. C is the maximum value of α , and determines the size of the margin. **ADD SOMETHING ABOUT WHICH WAY THAT WORKS HERE.** To return from α to the more familiar θ and θ_0 , we plug in:

$$\theta = \sum_{i=1}^n \alpha_i y^{(i)} x^{(i)} \quad (8)$$

$$\theta_0 = \frac{1}{\mathcal{M}} \left[\sum_{j \in \mathcal{M}} \left(y^{(j)} - \sum_{i \in \mathcal{S}} \alpha_i y^{(i)} (x^{(j)})^\top x^{(i)} \right) \right] \quad (9)$$

Where $\mathcal{M} = \{i : 0 < \alpha_i < C\}$ and $\mathcal{S} = \{i : 0 < \alpha_i\}$.

For example, given the data $X = [1, 2; 2, 2; 0, 0; -2, 3], Y = [1; 1; -1; -1]$, we generate the following objective function:

$$\min_{\alpha \in \mathbb{R}^n} \left[\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} (x^{(i)})^\top x^{(j)} - \sum_{i=1}^n \alpha_i \right] \quad (10)$$

This can be reformulated to be plugged into a standard quadratic programming solver, with **blah blah**




Figure 1: $C = 0.01, \sigma = 1$




Figure 2: $C = 0.1, \sigma = 1$




Figure 3: $C = 1, \sigma = 1$

We also note that as C increases, the number of boundary points decreases. Thus, the number of support vectors decrease. For the nonseparable example, for $C = (.01, .1, 1, 10, 100)$, we see that $\mathcal{M} = (400, 324, 133, 68, 59)$.

As per the usual procedure with the approximately equivalent λ , to select the correct C , we first compute α for various values of C using a training set, then determine the validation error using a validation set. We select the C which produces the lowest validation error in order to maximize the generality of our solution.

3 Titanic Data




Figure 4: $C = 10, \sigma = 1$




Figure 5: $C = 100, \sigma = 1$